


Quality Score — seed-level distribution (n = 10 seeds × datasets)



Each point is one generator-training seed for one dataset (same leakage-safe split).